



# Part 4: Transaction Management

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Transaction

- Example: Transfer of money from one account to another is a transaction consisting of two updates, one to each account
- A collection of operations that form a single logical unit of work
  - Transaction to **transfer \$50 from account A to account B**
    1. **read**( $A$ )
    2.  $A := A - 50$
    3. **write**( $A$ )
    4. **read**( $B$ )
    5.  $B := B + 50$
    6. **write**( $B$ )
  - If  $A = \$1000$  and  $B = \$2000$
  - If Success and only T1, *then at the end  $A = \$950$  and  $B = \$2050$ ,*
  - *Thus  $A + B = 3000$  (**before**  $1000 + 2000$  and **after**  $950 + 2050$ )*



# Issues

- In case of some **failure** in complete execution, partial effects of each **incomplete transaction be undone**, called **Atomicity**
- If **successful**, then its effects must persist in the database, called **Durability**
- **Multiple transactions** executing concurrently which may result in erroneous updates, require **Isolation**, to isolate transactions from others
- Chapter 14
  - Study of INDIVIDUAL transaction in detail
  - ACID Properties
  - Isolation by Serializability
- Chapter 15
  - Concurrency control techniques for **CONCURRENT** transactions Isolation
- Chapter 16
  - Recovery management for Atomicity and Durability for **CONCURRENT** transactions



# Chapter 14: Transactions

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use



# Chapter 14: Transactions

- INDIVIDUAL Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Transaction Definition in SQL
- Testing for Serializability



# Transaction Concept

- **Transaction**

- A *unit* of program execution that accesses and possibly updates various data items
- E.g. transaction to transfer \$50 from account A to account B:
  1. **read**( $A$ )
  2.  $A := A - 50$
  3. **write**( $A$ )
  4. **read**( $B$ )
  5.  $B := B + 50$
  6. **write**( $B$ )
- Main issues to deal with:
  - Completely execute
  - Concurrent execution of multiple transactions
  - Failures of various kinds, such as hardware failures and system crashes



# Example of Fund Transfer

- Transaction to transfer \$50 from account A to account B:
  1. **read**( $A$ )
  2.  $A := A - 50$
  3. **write**( $A$ )
  4. **read**( $B$ )
  5.  $B := B + 50$
  6. **write**( $B$ )

- **Suppose either of following happens**

1. **read**( $A$ )
2.  $A := A - 50$
3. **write**( $A$ )

**OR**

1. **read**( $A$ )
2.  $A := A - 50$
3. **write**( $A$ )
4. **read**( $B$ )
5.  $B := B + 50$



# Example of Fund Transfer

- **Atomicity requirement**
  - If the transaction FAILS after step 3 or before step 6,
    - 4 Failure could be due to software or hardware
    - 4 Money will be “LOST” leading to an INCONSISTENT database state
  - The system should ensure that updates of a partially executed transaction are not reflected in the database
  - In an atomic transaction,
    - 4 A series of database operations either ALL occur,
    - OR
    - 4 NOTHING occurs





# Example of Fund Transfer (Cont.)

- **Consistency requirement**
  - The information in the database must be kept in a consistent state
  - For example:
    - 4 If we have one table *B* that includes keys from another table *A* (such as **bank account transactions** that have account IDs referencing rows in a table of **bank accounts**),
      - Then **shouldn't allow deleting a row from *A* without deleting those rows from *B*** referring to the account being deleted
  - To ensure consistency for an individual transaction is the responsibility of the **application programmer** who codes the transaction



# Example of Fund Transfer (Cont.)

- **Consistency requirement**
  - Task facilitated by **automatic testing of integrity constraints**
  - **Explicitly specified integrity constraints** such as
    - 4 Primary keys and Foreign keys
  - **Implicit integrity constraints**
    - 4 e.g. sum of balances of all accounts, minus sum of loan amounts must equal value of cash-in-hand
  - In given example, The sum of A and B is unchanged by the execution of the transaction ( $A+B=3000$  (*before 1000+2000 and after 950+2050*))



# Example of Fund Transfer (Cont.)

- **Consistency requirement**
  - A transaction must see a consistent database
  - During transaction execution the database may be temporarily inconsistent
  - When the transaction completes successfully the database must be consistent
- 4 Erroneous transaction logic can lead to inconsistency



# Example of Fund Transfer (Cont.)

- Isolation requirement

**T1**

1. **read( $A$ )**
2.  $A := A - 50$
3. **write( $A$ )**

**T2**

4. **read( $B$ )**
5.  $B := B + 50$
6. **write( $B$ )**

**read( $A$ ), read( $B$ ), print( $A+B$ )**

- Computes  $A+B$

- *If  $A = \$1000$  and  $B = \$2000$*
- *And only T1, then at the end  $A = \$950$  and  $B = \$2050$ ,*
- *Thus  $A+B = 3000$  (before  $1000+2000$  and after  $950+2050$ )*

- **But, if T2 occurs, then T2 will observe an Inconsistent value**

- As it accesses the partially updated database, it will see an INCONSISTENT database (**the sum  $A + B = \$950 + \$2000 = \$2950$  will be less than it should be**)



# Example of Fund Transfer (Cont.)

- **Isolation requirement**

- If between steps 3 and 6, **another transaction T2 is allowed** to access the partially updated database, it will see an **INCONSISTENT** database (the sum  $A + B$  will be less than it should be)

**T1**

**T2**

1. **read( $A$ )**
2.  $A := A - 50$
3. **write( $A$ )**
4. **read( $B$ )**
5.  $B := B + 50$
6. **write( $B$ )**

**read( $A$ ), read( $B$ ), print( $A+B$ )**

- DB is temporarily inconsistent after Step 3, *If T2 concurrently occurs and reads  $A$  and  $B$  at intermediate point*
  - Computes  $A+B$ , *It will observe an Inconsistent value (2950 not 3000)*
  - **Furthermore, if T2 then performs updates on  $A$  and  $B$  based on the inconsistent values that it read, the database may be left in an inconsistent state even after both transactions have completed**



# Example of Fund Transfer (Cont.)

- **Isolation requirement**
  - Isolation can be ensured trivially by running transactions **serially**
    - 4 That is, one after the other
  - However, executing multiple transactions concurrently has significant benefits, as we will see later



# Example of Fund Transfer

- **Durability requirement**
  - Once the user has been notified that the transaction has completed (i.e., the transfer of the \$50 has taken place),
    - 4 Then the updates to the database by the transaction must **PERSIST**
      - Even if there are DBMS Software crashes or HARDWARE FAILURES



# ACID Properties

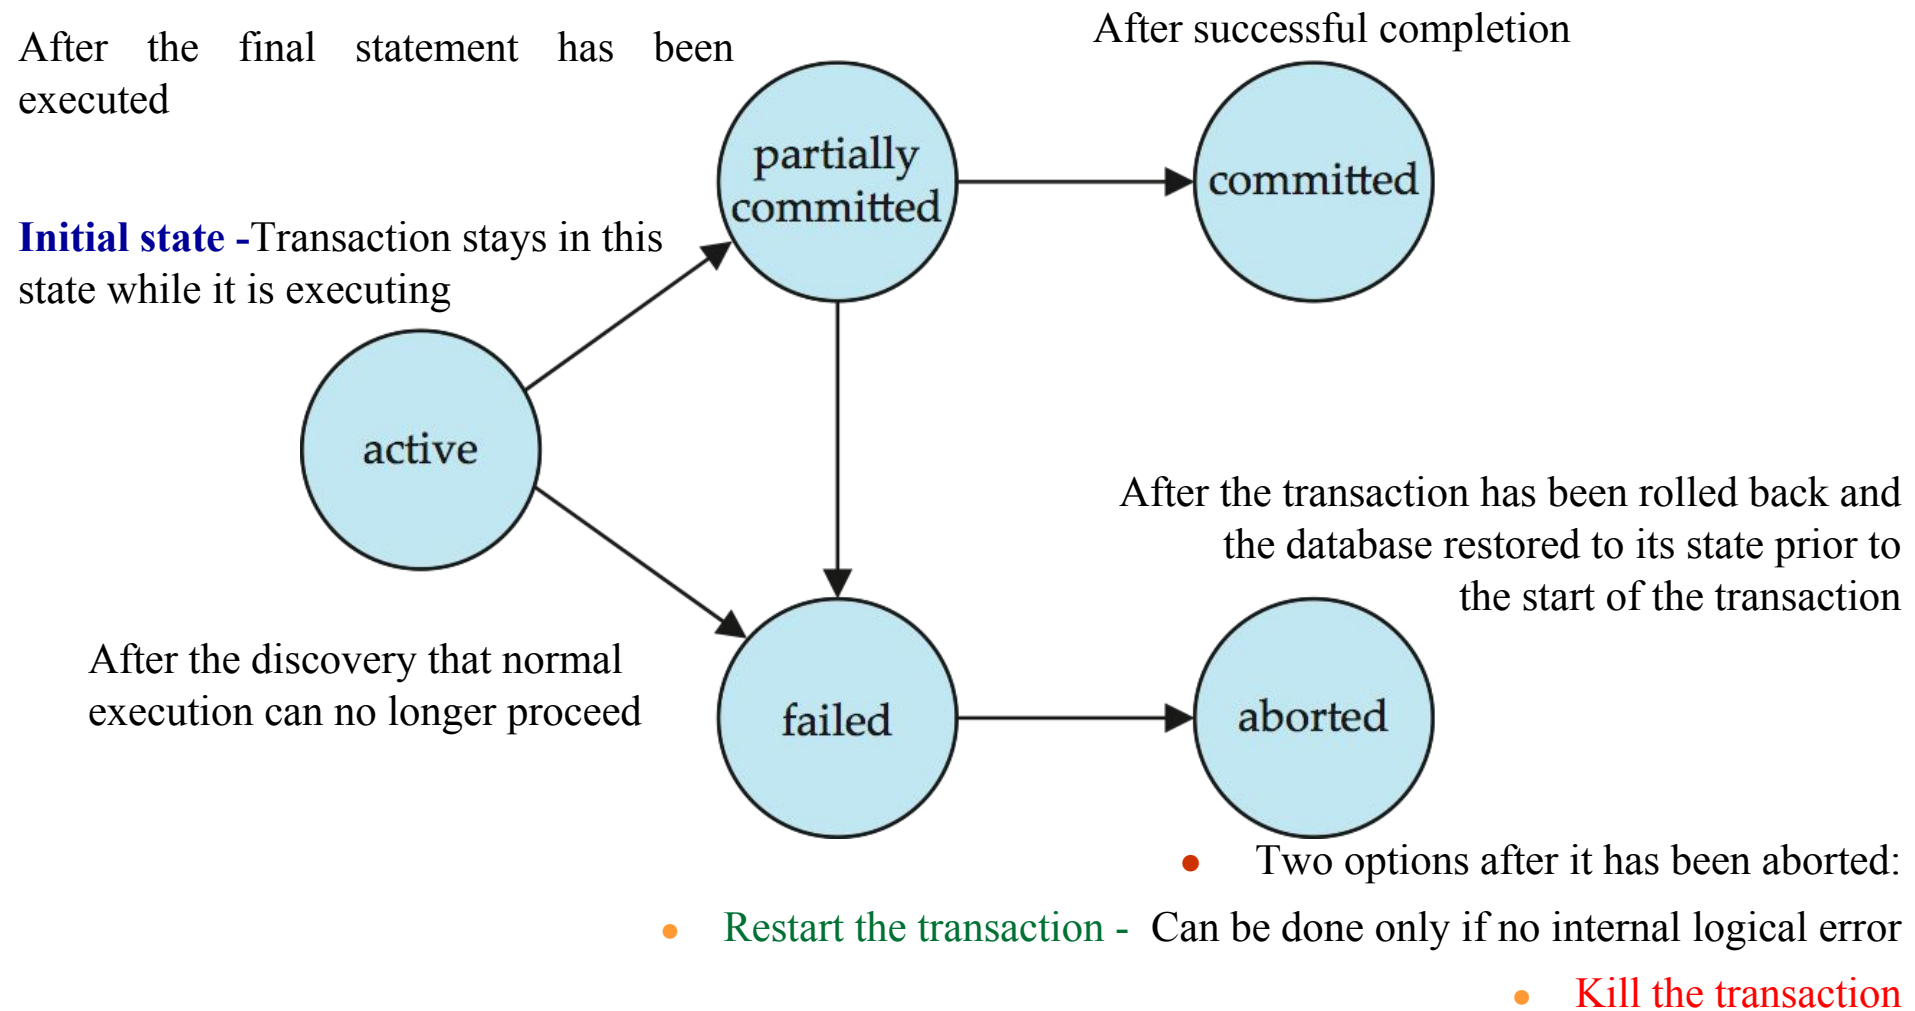
A **transaction** is a unit of program execution that accesses and possibly updates various data items. To preserve the integrity of data the database system must ensure:

- **Atomicity.** Either all operations of the transaction are properly reflected in the database or none are
- **Consistency.** Execution of a transaction in isolation preserves the consistency of the database
- **Isolation.** Although multiple transactions may execute concurrently, each transaction must be unaware of other concurrently executing transactions
  - Intermediate transaction results must be hidden from other concurrently executed transactions
    - 4 That is, for every pair of transactions  $T_i$  and  $T_j$ , it appears to  $T_i$  that either  $T_j$  finished execution before  $T_i$  started, or  $T_j$  started execution after  $T_i$  finished
- **Durability.** After a transaction completes successfully, the changes it has made to the database persist, even if there are system failures





# Transaction State (Cont.)





# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system
- Advantages
  - **Increased processor and disk utilization**, leading to better transaction *throughput*
    - 4 E.g. one transaction can be using the CPU while another is reading from or writing to the disk
  - **Reduced average response time** for transactions: short transactions need not wait behind long ones
- **Concurrency control schemes** – mechanisms to achieve isolation
  - To control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database
    - 4 Will study in Chapter 16, after studying notion of correctness of concurrent executions



# Schedules

- **Schedule**
  - A chronological execution sequence of transaction
    - 4 A schedule can have many transactions in it,
      - Each transaction comprising of number of instructions/tasks
  - **Must preserve the order** in which the **instructions appear in each individual transaction**
- It is some interleaving of the operations from the two transactions (without violating the order of operations within any individual transaction)
- A transaction that successfully completes its execution will have a commit instructions as the last statement (By default)
- A transaction that fails to successfully complete its execution will have an abort instruction as the last statement



# Example of Schedule

- Let
  - $T_1$  transfer \$50 from  $A$  to  $B$ , and
  - $T_2$  transfer 10% of the balance from  $A$  to  $B$
- Write the possible schedules.
- Verify:  $A=\$1000$ ,  $B=\$2000$ 
  - Before execution of both  $T_1$  and  $T_2$   $\square A+B=3000$
  - After execution of both  $T_1$  and  $T_2$   $\square A+B=3000$
- First, write operations:

Transaction 1:

$r1(A)$ ,  $A=A-50$ ,  $w1(A)$ ,  $r1(B)$ ,  $B=B+50$ ,  $w1(B)$ ,

Transaction 2:

$r2(A)$ ,  $tmp=A*0.10$ ,  $A=A-temp$ ,  $w2(A)$ ,  $r2(B)$ ,  $B=B+temp$ ,  $w2(B)$



# Schedule 1

- Let  $T_1$  transfer \$50 from  $A$  to  $B$ , and  $T_2$  transfer 10% of the balance from  $A$  to  $B$ .
- A **serial** schedule in which  $T_1$  is followed by  $T_2$  :

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

Suppose current values of accounts  
 $A = \$1000$  and  $B = \$2000$

**Before starting:  $A+B=\$3000$**

*Schedule 1:  $T_1$  followed by  $T_2$*

**After  $T_1$ :  $A=\$950$ ,  $B=\$2050$**

**After  $T_2$ :  $A = \$855(=950-95)$ ,  
 $B = \$2145(=2050+95)$**

**$A + B = \$855 + \$2145 = \$3000$**

*Preserved* after the execution of both transactions



## Schedule 2

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Suppose current values of accounts

$A = \$1000$  and  $B = \$2000$

**Before starting:  $A+B=\$3000$**

*Schedule 1:  $T_2$  followed by  $T_1$*

**After  $T_2$ :  $A=\$900(=1000-100)$ ,  
 $B=\$2100(=2000+100)$**

**After  $T_1$ :  $A = \$850, B = \$2150$**

**$A + B = \$850 + \$2150 = \$3000$**

*Preserved* after the execution of both transactions



# Serial Schedule

- $T_1$  and  $T_2$  schedules are **serial**:
  - Each serial schedule consists of a sequence of instructions from various transactions, where the instructions belonging to one single transaction appear together in that schedule
  - Recalling a well-known formula from combinatorics,
    - 4 For a set of  $n$  transactions, there exist  $n$  factorial ( $n!$ ) different valid serial schedules
- To describe consistent transaction behavior when transactions run at the same time uses transaction isolation model called **serializability**
- The serializable mode of transaction behavior tries to ensure that transactions run in such a way that they appear to be executed one at a time, or serially, rather than concurrently



# Schedule 3

- Let  $T_1$  and  $T_2$  be the transactions defined previously
- Following schedule is **NOT** a serial schedule, **but it is equivalent to Schedule 1**

$T_1$	$T_2$
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Current values of accounts

$A = \$1000$  and  $B = \$2000$

**Before starting:  $A+B=\$3000$**

**After Part1 of  $T_1$ :  $A=950, B=2000$**

**After Part 1 of  $T_2$ :  $A = 855(=950-95)$ ,  
 $B = 2000, temp=95$**

**After Part2 of  $T_1$ :  $A=855, B=2050$**

**After Part 2 of  $T_2$ :  $A = 855$ ,  
 $B = 2145(=2050+95)$**

**$A + B = \$855 + \$2145 = \$3000$**

*Preserved* after the execution of both transactions





# Schedule 4

Check the following concurrent schedule, is it preserving the value of  $(A + B)$  ?

$T_1$	$T_2$
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Suppose the current values of accounts  
 $A = \$1000$  and  $B = \$2000$

$T1: A=950$

$T2: A=1000$  (not written by T1)

$temp=100$

$A=900$

$B=2000$

$T1: A=950$

$B=2000$

$B=2050$

$B=2000+100=2100$

Final values after execution  $A = \$950, B = \$2100$

Total amount of money in accounts  $A$  and  $B$ , the sum  $A + B$  ( $= 950 + 2100 = 3050$ ) (Before execution  $= \$3000$ ) is NOT preserved after the execution of both transactions

Final state is an *inconsistent state*, since we have gained \$50

and the sum  $A + B$  is not preserved by the execution of the two transactions



# Serial Schedule

- When the database system **executes several transactions concurrently**, the corresponding schedule no longer needs to be serial
  - If two transactions are running concurrently, the **OS** may execute one transaction for a little while, then perform a context switch, execute the second transaction for some time, and then switch back to the first transaction for some time, and so on
  - With multiple transactions, the CPU time is shared among all the transactions
  - In general, it is **not possible to predict exactly how many instructions of a transaction will be executed** before the CPU switches to another transaction



# Serial Schedule

- If control of concurrent execution is left entirely to the **OS**, many possible schedules, including ones that leave the database in an inconsistent state
- **It is the job of the database system to ensure that any schedule that is executed will leave the database in a consistent state**
- **The concurrency-control component of the database system carries out this task**



# Serializability

- **Basic Assumption** – Each transaction preserves database consistency
  - Thus serial execution of a set of transactions preserves database consistency
  - A (possibly concurrent) schedule is **SERIALIZABLE**, if it is equivalent to a serial schedule
  - Different forms of schedule equivalence give rise to the notions of
1. **View serializability (Later Ch. 15 Concurrency)**
  2. **Conflict serializability**



# Serializability

## Conflict serializability

- Schedule is defined by equivalence to a serial schedule (no overlapping transactions) with the same transactions, **such that both schedules have the same sets of respective chronologically ordered pairs of conflicting operations (same precedence relations of respective conflicting operations)**
- Widely utilized
  - 4 Because it is easier to determine and covers a substantial portion of the view-serializable schedules



# *Simplified view of transactions*

- Since transactions are programs, it is difficult to determine exactly
  - What operations a transaction performs and
  - How operations of various transactions interact
- For this reason, here NOT considered the various types of operations that a transaction can perform on a data item,
  - Ignore operations other than **read** and **write** instructions
- Assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes
- Our simplified schedules consist of only **read** and **write** instructions
- **Commit operations, though relevant, are not considered until Section 14.7**



# Conflicting Instructions

- Instructions  $l_i$  and  $l_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if there exists some item  $Q$  accessed by both  $l_i$  and  $l_j$ , and at least one of these instructions wrote  $Q$ 
  1.  $l_i = \mathbf{read}(Q)$ ,  $l_j = \mathbf{read}(Q)$ .  $l_i$  and  $l_j$  don't conflict.
  2.  $l_i = \mathbf{read}(Q)$ ,  $l_j = \mathbf{write}(Q)$ . They conflict.
  3.  $l_i = \mathbf{write}(Q)$ ,  $l_j = \mathbf{read}(Q)$ . They conflict
  4.  $l_i = \mathbf{write}(Q)$ ,  $l_j = \mathbf{write}(Q)$ . They conflict
- Intuitively, a conflict between  $l_i$  and  $l_j$  forces a (logical) temporal order between them
  - If  $l_i$  and  $l_j$  are consecutive in a schedule and they do not conflict, their results would remain the same even if they had been interchanged in the schedule



# Conflict Serializability

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of non-conflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**
- A schedule  $S$  is **conflict serializable** if it is conflict equivalent to a serial schedule





## Conflict Serializability (Cont.)

- Check Schedule 3, is it a conflict serializable schedule.
- Sol: Try the possibility of swaps of non-conflicting instructions, to produce serial schedule.

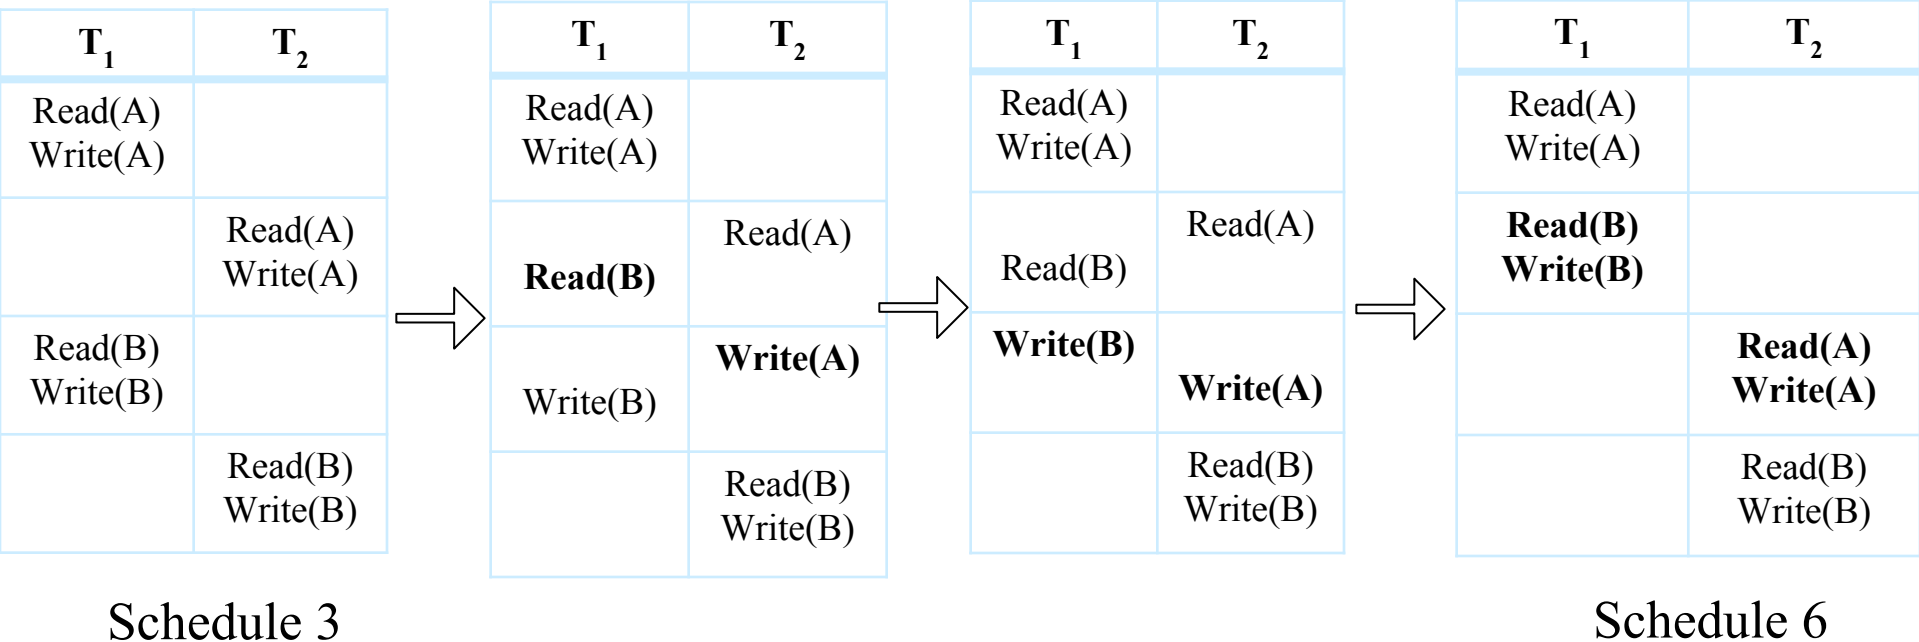
$T_1$	$T_2$
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

Schedule 3



# Conflict Serializability (Cont.)

- Check Schedule 3, is it a conflict serializable schedule.
- Sol: Try the possibility of swaps of non-conflicting instructions, to produce serial schedule.





# Conflict Serializability (Cont.)

- Schedule 3 can be transformed into Schedule 6, a serial schedule where  $T_2$  follows  $T_1$ , by series of swaps of non-conflicting instructions
- Therefore Schedule 3 is CONFLICT SERIALIZABLE**

$T_1$	$T_2$
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

Schedule 3

$T_1$	$T_2$
read (A) write (A) read (B) write (B)	read (A) write (A) read (B) write (B)

Schedule 6



# Conflict Serializability (Cont.)

- Check the serializability of following schedule.

$T_3$	$T_4$
read (Q)	write (Q)
write (Q)	

- Schedule is **not conflict serializable**
  - Unable to swap instructions in the above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$ , or the serial schedule  $\langle T_4, T_3 \rangle$
  - As read and write both are on single item Q



# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph**
  - Simple and Efficient method to determine conflict serializability of schedule
  - A directed graph where the Vertices are the transactions (names)
  - Draw an Edge(arc) from  $T_i$  to  $T_j$ 
    - 4 If the two transaction conflict, and  $T_i$  accessed the data item on which the conflict arose earlier
    - 4 Label the arc by the item that was accessed
    - 4 Draw edge from  $T_i \rightarrow T_j$  with label  $Q$  for one of three conditions:  

$T_i$



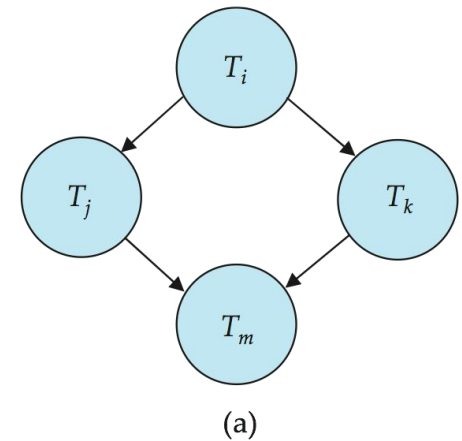
$T_j$

      1.  $write(Q) \quad read(Q)$
      2.  $read(Q) \quad write(Q)$
      3.  $write(Q) \quad write(Q)$



# Test for Conflict Serializability

- A schedule is **conflict serializable**
  - If and only if precedence graph is **ACYCLIC**
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph (Better algorithms take order  $n + e$  where  $e$  is the number of edges)
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph
  - This is a linear order consistent with the partial order of the graph





# Testing for Serializability

- Draw edge from  $T_i \rightarrow T_j$  for one of three conditions:



$T_i$                    $T_j$

1.  $write(Q)$      $read(Q)$
2.  $read(Q)$      $write(Q)$
3.  $write(Q)$      $write(Q)$

- **Example:      Schedule 1**

$T_1$	$T_2$
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

As all the instructions of  $T_1$  are executed before the  $T_2$



**CONFLICT SERIALIZABLE Schedule**





- $$T_i \quad T_j$$

- #### 4. Example: Schedule 2

$T_1$	$T_2$
<pre> read (A) A := A - 50 write (A) read (B) B := B + 50 write (B) commit </pre>	<pre> read (A) temp := A * 0.1 A := A - temp write (A) read (B) B := B + temp write (B) commit </pre>

**14.40**



## CONFLICT SERIALIZABLE Schedule



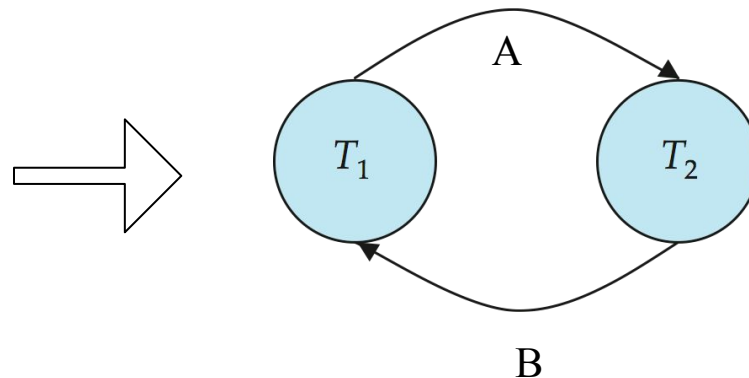


# Testing for Serializability

- **Example: Schedule 4**
- Contains the edge  $T_1 \rightarrow T_2$ , because  $T_1$  executes  $read(A)$  before  $T_2$  executes  $write(A)$
- Also contains the edge  $T_2 \rightarrow T_1$ , because  $T_2$  executes  $read(B)$  before  $T_1$  executes  $write(B)$

$T_1$	$T_2$
read (A) $A := A - 50$	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	
	$B := B + temp$ write (B) commit

**NOT CONFLICT SERIALIZABLE Schedule**



**Cycle is present in graph**



# Testing for Conflict Serializability

## Example: Schedule A

W3 (A)

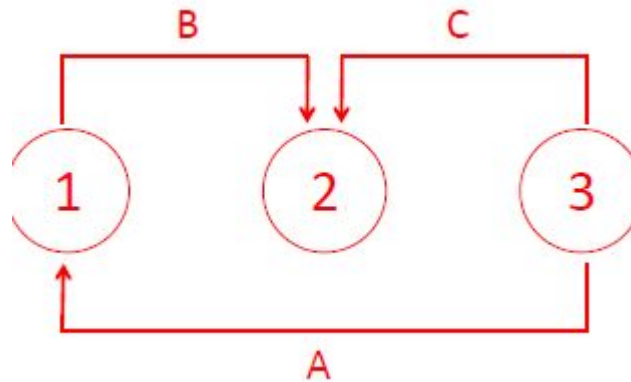
R1 (A)

W1 (B)

R2 (B)

W3 (C)

R2 (C)



No cycles: **YES, conflict serializable**




# Test for Conflict Serializability

## Example: Schedule A

W3 (A)  
R1 (A)  
W1 (B)  
R2 (B)  
W3(C)  
R2 (C)

## Serial Equivalent Schedule?

W3 (A)  
W3(C)  
R1 (A)  
W1 (B)  
R2 (B)  
R2 (C)



No cycles: **YES, conflict serializable**

**Only serial equivalent schedule: T3, T1, T2**

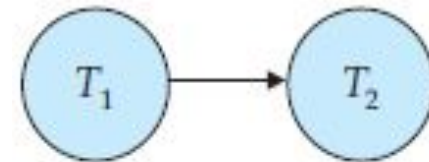
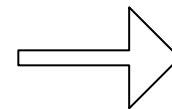


# Test for Conflict Serializability

- **Example: Schedule 3**

$T_1$	$T_2$
read (A) write (A)	read (A) write (A)
read (B) write (B)	read (B) write (B)

As all the instructions of  $T_1$  are executed before the completion of  $T_2$



**CONFLICT SERIALIZABLE Schedule**



# Transaction Failure

- Upto this considered that **NO TRANSACTION FAILURE**
  - But, this is not really true

$T_8$	$T_9$
read (A) write (A)	read (A) commit
read (B)	

- **IF TRANSACTION FAILURE** is there,
  - Need to **UNDO the effect of that transaction** to ensure **ATOMICITY**
  - Requires that any transaction  $T_j$  that dependent on  $T_i$  is also aborted
  - Achieved by restricted type of schedules

4 Recoverable Schedule and Cascadeless Schedule



# Recoverable Schedules

- **Recoverable schedule** — if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$

- The Schedule 11

$T_8$	$T_9$
read (A)	
write (A)	
	read (A)
	commit
read (B)	

- Only commit is with  $T_9$  not with  $T_8$  (**Partial Schedule**)
- $T_9$  reads the value written by  $T_8$  ( **$T_9$  is dependent upon  $T_8$** )
- If  $T_8$  is failed, Requires Undo for  $T_9$  also, which is not possible



# Recoverable Schedules

Schedule 11

$T_8$	$T_9$
read (A) write (A)	
	read (A) commit
read (B)	

- Database must ensure that schedules are recoverable
  - $T_9$  is **dependent** on  $T_8$ 
    - Because of this, must abort  $T_9$  to ensure atomicity
    - However,  $T_9$  has already committed and cannot be aborted
      - » Thus, a situation where it is impossible to recover correctly from the failure of  $T_8$
  - If  $T_8$  should abort,  $T_9$  would have read (and possibly shown to the user) an inconsistent database state
  - Schedule is **NOT RECOVERABLE**



# Recoverable Schedules

Schedule 10

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)

- Database must ensure that schedules are recoverable
  - $T_{11}$  is dependent upon  $T_{10}$  and  $T_{12}$  is dependent on  $T_{11}$ 
    - Abort is there with  $T_{10}$ , it is possible to abort  $T_{11}$  and  $T_{12}$
  - Schedule is **RECOVERABLE**





# Cascading Rollbacks

- The schedule is **RECOVERABLE, BUT REQUIRES CASCADING ROLLBACKS**
  - If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back
- Can lead to the undoing of a significant amount of work

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	read (A)
abort		

- A single transaction failure leads to a series of transaction rollbacks
  - As none of the transactions has committed till abort of  $T_{10}$



# Cascadeless Schedules

- Cascading rollback is undesirable
  - As requires a lot many rollback
- For each pair of transactions  $T_x$  and  $T_y$  such that  $T_y$  reads a data item previously written by  $T_x$ , the commit operation of  $T_x$  appears before the read operation of  $T_y$

$T_x$        $T_y$

Write(Q)

Commit

Read(Q)

- Every CASCADELESS schedule is RECOVERABLE schedule



# Concurrency Control

- If only one transaction is allowed to execute at a time
  - It generates serial schedules
  - BUT, provides a poor degree of concurrency
- For concurrent transactions, DBMS need to check
  - Either conflict or view serializable schedule ?
  - Recoverable and preferably cascadeless schedule ?
  - Are serial schedules recoverable/cascadeless ?



# Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency, allowing schedules that are not serializable
  - E.g. a read-only transaction that wants to get an approximate total balance of all accounts
  - E.g. database statistics computed for query optimization can be approximate (why?)
  - Such transactions need not be serializable with respect to other transactions
- Tradeoff accuracy for performance
  - Allowing not to wait for other transactions serial order
    - 4 E.g. Permit the transaction to read a data item even if it was written by a transaction that has not been committed
  - To benefit the long transactions whose results do not need to be precise



# Transaction Isolation Levels of Consistency

- With different possible outcomes for the same transaction scenario
  - The same work performed in the same fashion with the same inputs may result in different answers, depending on your isolation level
  - Isolation levels
    - 4 **Serializable (Default)**
    - 4 **Repeatable read**
    - 4 **Read committed**
    - 4 **Read uncommitted**
  - These levels are defined in terms of three phenomena that are either permitted or not permitted at a given isolation level



# Transaction Isolation Levels of Consistency

- **Serializable (Default)**
  - Serializable schedule are the ideal way to ensure consistency, but in our day-to-day lives, NOT IMPOSING SUCH STRINGENT REQUIREMENT
  - Ensured serializability only for the transactions that run on the database, without user interaction
    - In Online Shopping, when user surfing for an item it is available in stock, but by the time user goes through the checkout process, that item might no longer be available □ this would be a nonrepeatable read
- **Repeatable read**
- **Read committed**
- **Read uncommitted**



# Transaction Isolation Levels of Consistency

- **Serializable (Default)**
  - Ensured serializability only for the transactions that run on the database, without user interaction
- **Repeatable read**
  - **Allows only committed data to be read** and further requires that, between two reads of a data item by a transaction, **no other transaction is allowed to update it**
  - Only committed records to be read, repeated reads of same record must **return same value**
  - However, a transaction may not be serializable with respect to other transactions, so it may find some records inserted by a transaction but not find others
- 4 For instance, when it is searching for data satisfying some conditions, a transaction may find some of the data inserted by a committed transaction, but may not find other data inserted by the same transaction (until it is committed-as this will read only committed one)
- **Read committed**
- **Read uncommitted**



# Transaction Isolation Levels of Consistency

- **Serializable (Default)**
- **Repeatable read**
- **Read committed**
  - Allows only committed data to be read, but successive reads of record may return different (but committed) values
    - 4 Does not require repeatable reads
  - For instance, between two reads of a data item by the transaction, another transaction **may have updated the data item and committed**
- **Read uncommitted**
  - Allows uncommitted data to be read and the lowest isolation level allowed by SQL





# Transaction Isolation Levels of Consistency

## Read phenomena that violate serializability

- Three different *read phenomena* when Transaction 1 reads data that Transaction 2 might have changed
  1. **Dirty read**
  2. **Nonrepeatable read**
  3. **Phantom read**
- Examples, two transactions queries use the data table: USERS

id	name	age
1	Joe	20
2	Jill	25

1. First Query 1 is performed
2. Then, in the second transaction, Query 2 is performed on same table
3. Finally, in the first transaction, Query 1 is performed again



# Transaction Isolation Levels of Consistency

## Read phenomena - Dirty read - Aka *uncommitted dependency*

- Occurs when a transaction is allowed to read data from a row that has been **modified by another** running transaction and **not yet committed**

id	name	age
1	Joe	20
2	Jill	25

Transaction 1

```
/* Query 1 */ SELECT age FROM users  
WHERE id = 1; /* will read 20 */
```

```
/* Query 1 */ SELECT age FROM users  
WHERE id = 1;
```

Reads the uncommitted data  
*will read 21 for id=1*

Changes (already read by Transaction 1) or updates different changes to the database, then the view of the data may be wrong in the records of Transaction 1

Transaction 2

Changes a row,  
but does not commit the changes

```
/* Query 2 */ UPDATE users SET age = 21  
WHERE id = 1; /* No commit here */
```

**ROLLBACK;**



# Transaction Isolation Levels of Consistency

- **Read phenomena - Non-repeatable reads**
  - Occurs, when during the course of a transaction, a row is retrieved twice and the values within the row differ between reads

id	name	age
1	Joe	20
2	Jill	25

Transaction 1

```
/* Query 1 */ SELECT age FROM users  
WHERE id = 1; /* will read 20 */
```

```
/* Query 1 */ SELECT age FROM users  
WHERE id = 1; will read 21  
COMMIT;
```

Transaction 2

Changes a row, with commit

```
/* Query 2 */ UPDATE users SET age = 21  
WHERE id = 1; /* No commit here */  
COMMIT;
```

Transaction 1 has already seen a different value for *age* in that row – as non-repeatable reads



# Transaction Isolation Levels of Consistency

- **Read phenomena - Phantom reads**
  - Occurs when, in the course of a transaction, two identical queries are executed, and the collection of rows returned by the second query is different from the first
  - A special case of *Non-repeatable reads* when Transaction 1 repeats a ranged *SELECT ... WHERE* query and, between both operations, Transaction 2 creates (i.e. INSERT) new rows (in the target table) which fulfill that *WHERE* clause

		id	name	age
Transaction 1		1	Joe	20
Transaction 2		2	Jill	25

Transaction 1		/* Query 1 */ <b>SELECT * FROM</b> users <b>WHERE</b> age <b>BETWEEN</b> 10 <b>AND</b> 30;		
Transaction 2		/* Query 2 */ <b>INSERT INTO</b> users(id,name,age) <b>VALUES</b> ( 3, 'Bob', 27 ); <b>COMMIT</b> ;		
Transaction 1		/* Query 1 */ <b>SELECT * FROM</b> users <b>WHERE</b> age <b>BETWEEN</b> 10 <b>AND</b> 30; <b>COMMIT</b> ;		

A different set of rows may be returned the second time



# Transaction Isolation Levels of Consistency

- **Dirty read**
  - The meaning of this term is as bad as it sounds
  - Related to reading UNCOMMITTED data, and occur when an UPDATE, INSERT, or DELETE from another transaction is read, and the other transaction has NOT yet committed the data
  - Permitted to read uncommitted or dirty data
  - It is reading "in progress" data, which may not be complete, and may never actually be committed
  - Example: Open an OS file that someone else is writing and reading whatever data happens to be there
  - Data integrity is compromised, foreign keys are violated, and unique constraints are ignored



# Transaction Isolation Levels of Consistency

- **Non-repeatable and phantom read**
  - Deal with data modification operations from a different transaction, which were committed after your transaction began, and then read by your transaction
- **Nonrepeatable read**
  - Simply means that if you read a row at time T1 and try to reread that row at time T2, the row may have changed
  - Transaction reads committed **UPDATES** from another transaction
  - It may have disappeared, it may have been updated, and so on
  - The same row now has different values than it did when your transaction began



# Transaction Isolation Levels of Consistency

- **Non-repeatable and phantom read**
  - Deal with data modification operations from a different transaction, which were committed after your transaction began, and then read by your transaction
- **Phantom read**
  - Similar to non-repeatable, but when reading from committed **INSERTS** and/or **DELETES** from another transaction
  - Means that if you execute a query at time T1 and re-execute it at time T2, additional rows may have been added to the database, which may affect your results
  - Differs from a non-repeatable read in that with a phantom read, data you already read hasn't been changed, but instead, more data satisfies your query criteria than before
    - 4 There are new rows or rows that have disappeared since you began the transaction



# Transaction Isolation Levels of Consistency

## Isolation Levels vs Read Phenomena

Isolation level	Dirty reads	Non-repeatable reads	Phantom reads
Read Uncommitted	may occur	may occur	may occur
Read Committed	-	may occur	may occur
Repeatable Read	-	-	may occur
Serializable	-	-	-





# Transaction Isolation Levels of Consistency

- Explicitly you can set the isolation level
  - SET TRANSACTION ISOLATION LEVEL SERIALIZABLE
  - SET TRANSACTION ISOLATION LEVEL READ COMMITTED
  - SET TRANSACTION ISOLATION LEVEL REPEATABLE READ
  - SET TRANSACTION ISOLATION LEVEL READ UNCOMMITTED
  - Oracle supports first two only and one more READ ONLY
  - To change, this setting has to be the first statement of a transaction
  - Also Turn OFF “Auto Commit”



# Concurrency Control

- Implementation of Isolation level
  - Upto this, seen properties of schedule to leave the database in consistent state and allow transaction failures to be handled in a safe manner
  - Concurrency control policies
    - 4 When multiple transactions executed concurrently, only **acceptable schedules are generated**, regardless of how the OS time shares resources (such as CPU time) among the transactions



# Concurrency Control

- Implementation of Isolation level
  - Upto this, seen properties of schedule to leave the database in consistent state and allow transaction failures to be handled in a safe manner
  - Concurrency control policies
    - 4 Example: Transaction acquires a lock on the entire database before it starts and releases the lock after it has committed
      - While a transaction holds a lock, no other transaction is allowed to acquire the lock, and all must therefore wait for the lock to be released
        - » Only one transaction can execute at a time, Serial schedule which is recoverable and cascadeless
        - » POOR PERFORMANCE, only one transaction, No concurrency



# Concurrency Control

- Implementation of Isolation level
  - To provide high degree of concurrency with assurance of generated schedules are conflict or view serializable, recoverable and cascadeless

4 Ch. 15 Concurrency Control



# Transaction Definition in SQL

- Data manipulation language must include a construct for specifying the set of actions that comprise a transaction
- In SQL, a transaction begins implicitly
- A transaction in SQL ends by:
  - **Commit work** commits current transaction and begins a new one
  - **Rollback work** causes current transaction to abort
- In almost all database systems, by default, every SQL statement also commits implicitly if it executes successfully
  - Implicit commit can be turned off by a database directive
    - 4 E.g. in JDBC, `connection.setAutoCommit(false);`



# ACID Properties

## To Achieve ACID Properties

- **Atomicity.** Recovery
- **Consistency.** Developer
- **Isolation.** Concurrency Controller
- **Durability.** Recovery



# End of Chapter 14

**Database System Concepts, 6<sup>th</sup> Ed.**

©Silberschatz, Korth and Sudarshan  
See [www.db-book.com](http://www.db-book.com) for conditions on re-use